

Evaluating OpenCode's Free Models

OpenCode Zen currently offers a set of free, time-limited promotional models. They are all temporary, and OpenCode is explicit that usage data may be logged or used for training during this period — worth keeping in mind for sensitive work.

What Each Model Actually Is

MiMo V2.5 Free From Xiaomi's MiMo team. Their V2-Flash technical report describes a large Mixture-of-Experts model (309B total parameters, 15B active) using a hybrid attention architecture that interleaves local sliding-window and global attention, built for efficient long-context reasoning and agentic tasks. Strong general-purpose and coding all-rounder, optimized for speed and cost efficiency.

Nemotron 3 Ultra Free NVIDIA's largest model in the new Nemotron 3 family (Nano / Super / Ultra). It uses a Mixture-of-Experts hybrid Mamba-Transformer architecture for high throughput and context lengths up to 1M tokens, with reinforcement-learning post-training for reasoning, multi-step tool use, and adjustable reasoning budgets. As the "Ultra" tier, it's the most capable and heaviest Nemotron 3 variant — a good pick if you need very long context or complex agentic tool-calling.

DeepSeek V4 Flash Free The fast, lightweight variant of DeepSeek's V4 line. "Flash" models trade some raw capability for speed and cost — good for quick iteration, less ideal for the hardest reasoning tasks.

North Mini Code Free A smaller, code-focused model. "Mini" suggests it's optimized for fast, cheap code completion and edits rather than complex multi-file reasoning — best for quick snippets, not architecture-level work.

Big Pickle This one's genuinely a mystery. OpenCode itself just calls it "a stealth model that's free on OpenCode for a limited time" and doesn't disclose what it actually is. It has a 200k context window with reasoning enabled, and some unverified online speculation ties it to GLM, but that's not confirmed. Anecdotal reports describe it as a strong, well-rounded coding performer — sometimes claimed to beat competing tools — but there's no independent, rigorous benchmark data backing that up, since the model's identity isn't even disclosed.

Recommendation by Use Case

Priority	Pick
Best raw reasoning / longest context / complex agentic workflows	Nemotron 3 Ultra
Best general-purpose balance of speed + capability	MiMo V2.5
Fast, cheap, lots of iteration	DeepSeek V4 Flash
Quick code completions only	North Mini Code
Curious/willing to gamble on an unknown but reportedly strong	Big Pickle

Bottom Line

For general coding work, start with **Nemotron 3 Ultra** for anything needing deep reasoning or huge context, and **MiMo V2.5** as a daily driver for typical tasks. Big Pickle is tempting because of the buzz, but since OpenCode won't even say what it is, it's best treated as an experiment rather than a reliable choice for anything important — and there are no independent, rigorous benchmarks yet for any of these, since they're brand-new free releases.

Caveat: Since these are all “free trial” models, expect them to either get pulled, repriced, or replaced soon — worth not building critical workflows around any of them long-term.